

DiagramIQ

It turns the way work is *described* into the way work is *modelled*.

In one sentence, for anyone: hand it a standard operating procedure, a meeting transcript, a spreadsheet, a Visio file or a photograph of a whiteboard — it reads the material the way a business analyst would, shows you what it understood so a human can correct it, and then produces a proper process map: the BPMN 2.0 kind that opens straight in Signavio, Camunda or bpmn.io.

80%

less analyst effort per
process map

6h → 1.2h

per process, end to end

~2min

to a reviewable first draft

6

input types accepted, no re-
keying

WHAT IT ACCEPTS

SOP / work instruction (Word, PDF)

Meeting transcript or notes

Photo or image of a diagram

Process spreadsheet

Visio drawing

Existing BPMN / XML

WHAT IT RETURNS

BPMN 2.0 XML + rendered diagram

Process discovery workbook

Conformance audit

The questions to ask the business

1

Upload

SOP, transcript, image,
spreadsheet or Visio

→

2

The AI reads it

words and embedded
diagrams, together

→

3

A human approves

the understanding —
before anything is drawn

→

4

BPMN 2.0 out

built, laid out and audited

WHO IT IS FOR

Transformation and consulting teams who must map a client's processes before they can improve them; automation teams who need a reliable "as-is" before designing the "to-be"; and risk and controls teams who need the documented process and the modelled one to be the same thing.

\$1.14m

illustrative annual delivery cost
released
at 25 engagements a year — see
page 5

01 The problem, in plain terms

Everybody has the knowledge written down. Almost nobody has it modelled.

A client's process already exists on paper — in an SOP, a policy, a training deck, a recorded workshop. Turning that into a process map that a consultant, an automation team or an auditor can actually use is done **by hand, one box at a time**. It is slow, it is expensive, and the person doing it is a scarce specialist.

Where six hours per process actually goes

Reading and understanding the source		1.2 h
Drafting the steps, roles and systems		1.5 h
Drawing the BPMN — lanes, gateways, layout		2.0 h
Fixing naming and modelling conventions		0.8 h
Review rework after the first pass		0.5 h
Total, per single process		6.0 h

Multiply it. A mid-sized transformation maps **250 processes**. That is **1,500 analyst hours** — roughly **9 months of one person's working year** — spent drawing boxes rather than improving the process inside them.

And the tools that exist do not close the gap

Modelling editors

Visio, Camunda Modeler, Signavio. Excellent drawing surfaces — but they read nothing. Every box is still placed by a human who has already read the document.

Process mining

Celonis, ABBYY, UiPath. Needs an event log from a transactional system. A process that only exists in a document or in someone's head produces no log at all.

Diagram / OCR converters

Recover shapes and text from a picture. A rectangle becomes a rectangle — not a task, owned by a role, conditional on a threshold. The meaning is left behind.

A chat AI, used directly

Produces XML that may not open, carries no layout, and gives the reviewer no way to see what it understood before accepting it — or to fix one step without regenerating everything.

What that costs, at the scale it is actually done

1,500_h

of specialist time per engagement,
at 250 processes

\$57,000


delivery cost per engagement,
at \$38 per hour blended

\$1.43_m

a year across 25 engagements —
spent drawing, not advising


02 The solution — how it works

Four steps. The third one is the reason it can be trusted.



Why step 3 changes everything. Every other AI tool asks you to check its *output* — a finished diagram, where an error is expensive to find and expensive to fix. DiagramIQ asks you to check its *understanding*, in a table, in thirty seconds, before a single box exists. Corrections cost one keystroke instead of one re-draw.


What you get back, every time



And it keeps working with you

Describe a correction in ordinary English — “*Approve request belongs to Finance, not Procurement*” — and it applies it, then tells you exactly what it changed and what it did not. If a revision makes things worse, **roll back** to the previous version. Export the steps to Excel, edit them in the tool everyone already has, and re-import: the edits **patch** the existing diagram instead of duplicating it.

A worked example — one real SOP, start to finish



Reference sections — purpose, scope, definitions, revision history — were correctly left out of the process. The approval thresholds written into the clause text were carried into the diagram as decision criteria.

03 What makes it different in the market

Three things no competing category does together.

CAPABILITY	MODELLING EDITORS	PROCESS MINING	OCR / DIAGRAM CONVERTERS	GENERAL CHAT AI	DIAGRAMIQ
Works with no event log / no system data	✓	×	✓	✓	✓
Reads a document and extracts the process	×	×	part	part	✓
Reads text <i>and</i> the diagram inside it, together	×	×	×	×	✓
Human approves the interpretation before drawing	×	×	×	×	✓
Guaranteed valid BPMN 2.0 with layout	✓	part	×	×	✓
Audits itself against modelling conventions	part	×	×	×	✓
Says what it could not read, and what to ask	×	×	×	×	✓
Edits round-trip through Excel without duplicating	×	×	×	×	✓

USP 1

It reads the picture and the paragraph as one

An SOP explains a rule in words on page 2 and draws it on page 4. Every other tool sees one or the other. DiagramIQ sends both to the model together, positioned where they sat in the document.

USP 2

The human signs off the understanding, not the drawing

Verification happens on a table of steps, before generation. And the diagram itself is built by fixed rules — so the AI's judgement is reviewed, and the AI never gets to write invalid XML.

USP 3

It is honest about what it does not know

Missing thresholds become questions for the business, not invented answers. A partial reading is labelled as partial. That is what makes the output safe to put in front of a client.

How to prove it on a live account — a 3-week pilot

WEEK	WHAT HAPPENS	WHAT IT PRODUCES
Week 1	Pick 20 processes already documented on the account. Time how long two of them take to model by hand today — that is the baseline, measured rather than assumed.	A real baseline in hours
Week 2	Run all 20 through DiagramIQ. Each one is reviewed and approved by the same analysts who would have drawn it. Record the actual review time per process.	20 diagrams + 20 timings
Week 3	Compare quality against the hand-drawn versions, count the audit findings and the questions raised for the business, and recompute page 5 with the measured numbers.	A defensible business case

04 The numbers — effort, cost and revenue

Every figure below is computed from the four assumptions in the box. Change them and the model moves with you.

Assumptions used

6.0 h
manual effort per process map

1.2 h
with DiagramIQ, incl. review

250
processes in a typical engagement

\$38
blended delivery cost per hour

Effort and cost released

SCOPE	MANUAL	WITH DIAGRAMIQ	HOURS SAVED	COST RELEASED
Per single process	6.0 h	1.2 h	4.8 h	\$182
Per 100 processes	600 h	120 h	480 h	\$18,240
Per engagement (250 processes)	1,500 h	300 h	1,200 h	\$45,600
Across 25 engagements a year	37,500 h	7,500 h	30,000 h	\$1.14 M

30,000_h

analyst capacity released each year — about 16 FTEs redeployed to billable work

\$1.14_M

annual delivery cost taken out at 25 engagements per year

8 → 2.5

weeks to complete discovery and modelling — a 69% shorter phase

How that becomes revenue and margin for Genpact

LEVER	WHAT CHANGES	ILLUSTRATIVE VALUE
Margin on fixed-price work	The saving lands straight in delivery cost. On a \$400k engagement carrying \$260k of delivery cost, removing \$45.6k lifts gross margin from 35% to 46%.	+11 pts
Capacity redeployed	30,000 hours of scarce BA time moved off box-drawing and onto advisory and automation design — billable work that does not need new hiring.	≈ 16 FTEs
Faster time to value	Discovery closes 5.5 weeks earlier, so the automation and re-engineering phase — the larger, higher-margin scope — starts and bills sooner.	–69% cycle
Win rate and pricing	Bid a discovery phase at a price a manual competitor cannot match, while keeping the margin, because the cost base is genuinely lower.	Competitive edge
A productised offering	"Process discovery as a service" — the same accelerator resold across accounts, with the marginal cost of one more process close to zero.	New revenue line
Combined effect	Lower cost to serve, more billable capacity, shorter path to the bigger scope.	\$1.14 M + upside